

Exercise 7 - Association Rules

First name: Brian

Last name: Schweigler

Matriculation number: 16-102-071

(1) Take the titanic dataset and using all attributes to predict the class 'Survived' (convert age and fare into classes ; exclude names from the attribute list) Build a Support vector machines (SVM) model with:

(a) Linear kernel (b) Polynomial kernel (c) radial basis function (RBF) kernel (d) sigmoid kernel Show the Comparison of the performances.

In [1]:

```
%load_ext autoreload
%autoreload 2
%matplotlib inline

import pandas as pd
import numpy as np
from sklearn import preprocessing
from mlxtend.evaluate import accuracy_score
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.svm import SVC
from sklearn.model_selection import KFold

df = pd.read_csv("data/titanic.csv", index_col='Name')
pd.set_option('display.max_colwidth', None)
le = preprocessing.LabelEncoder()
bins = [0, 4, 18, 65, 100]
labels = ['Infant', 'Child', 'Adult', 'Elderly']
labels = [1, 2, 3, 4]
df['AgeGroup'] = pd.cut(df['Age'], bins=bins, labels=labels, right=False)
df["AgeGroup"] = le.fit_transform(df["AgeGroup"])
df['FareGroup'] = pd.qcut(x=df['Fare'], q=4)
df["FareGroup"] = le.fit_transform(df["FareGroup"])
df["Survived"] = le.fit_transform(df["Survived"])
df["Sex"] = le.fit_transform(df["Sex"])
print(df.head(2))

models = {
    'Linear SVM Kernel': SVC(kernel='linear'),
    'Poly SVM Kernel': SVC(kernel='poly'),
    'Rbf SVM Kernel': SVC(kernel='rbf'),
    'Sigmoid SVM Kernel': SVC(kernel='sigmoid'),
}
```

	Survived	Pclass	Sex	\
Name				
Mr. Owen Harris Braund	0	3	1	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	1	1	0	

	Age	\
Name		
Mr. Owen Harris Braund	22.0	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	38.0	

	Siblings/Spouses Aboard	\
Name		
Mr. Owen Harris Braund	1	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	1	

	Parents/Children Aboard	\
Name		
Mr. Owen Harris Braund	0	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	0	

	Fare	AgeGroup	\
Name			
Mr. Owen Harris Braund	7.2500	2	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	71.2833	2	

	FareGroup
Name	

Name	0
Mr. Owen Harris Braund	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	3

In [2]:

```
all_features = ['Pclass', 'Sex', 'Siblings/Spouses Aboard',
               'Parents/Children Aboard', 'AgeGroup', 'FareGroup']

def kfold_svm_eval(model, x: pd.DataFrame, y: pd.DataFrame):
    kf = KFold(n_splits=5, shuffle=True, random_state=6)
    accuracy = np.empty(kf.n_splits)
    precision = np.empty(kf.n_splits)
    recall = np.empty(kf.n_splits)
    f1 = np.empty(kf.n_splits)

    i = 0
    for train_index, test_index in kf.split(x):
        x_train, x_test = x.iloc[train_index], x.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        model.fit(x_train, y_train)
        prediction = model.predict(x_test)

        accuracy[i] = accuracy_score(prediction, y_test)
        precision[i] = precision_score(prediction, y_test)
        recall[i] = recall_score(prediction, y_test)
        f1[i] = f1_score(prediction, y_test)
        i += 1

    print(name)
    return accuracy, precision, recall, f1

for name, model in models.items():
    a, p, r, f = kfold_svm_eval(model, df[all_features], df["Survived"])
    data = {
        'Fold': [1, 2, 3, 4, 5],
        'Accuracy': a,
        'Precision': p,
        'Recall': r,
        'F1-Score': f,
    }

    scores = pd.DataFrame(data).set_index('Fold')
    display(scores)
```

Linear SVM Kernel

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.842697	0.705882	0.857143	0.774194
2	0.758427	0.625000	0.737705	0.676692
3	0.762712	0.701493	0.681159	0.691176
4	0.745763	0.685714	0.676056	0.680851
5	0.819209	0.692308	0.789474	0.737705

Poly SVM Kernel

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.848315	0.735294	0.847458	0.787402
2	0.792135	0.638889	0.807018	0.713178
3	0.785311	0.761194	0.698630	0.728571
4	0.768362	0.685714	0.716418	0.700730
5	0.841808	0.723077	0.824561	0.770492

Rbf SVM Kernel

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.865169	0.720588	0.907407	0.803279
2	0.803371	0.638889	0.836364	0.724409
3	0.796610	0.731343	0.731343	0.731343
4	0.774011	0.685714	0.727273	0.705882
5	0.853107	0.738462	0.842105	0.786885

Sigmoid SVM Kernel

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.539326	0.338235	0.383333	0.359375
2	0.589888	0.388889	0.491228	0.434109
3	0.587571	0.447761	0.454545	0.451128
4	0.502825	0.400000	0.378378	0.388889
5	0.468927	0.307692	0.289855	0.298507

Seems like the linear and polynomial models have the highest accuracy, while the Sigmoid model performs the worst. We expect a clear correlation between e.g. PClass and Survived, thus a Sigmoid shouldn't fit, which is confirmed.

Build Support vector machines (SVM) model with your selected stock / market index using all attributes to predict 'daily returns'(decision). ('daily returns' must first be converted into a decision class that will be used as the target(label), all other attributes must be grouped into classes)

Explain how the following kernels affects the performance of the model. (a) Linear kernel (b) Polynomial kernel (c) radial basis function (RBF) kernel (d) sigmoid kernel Show the Comparison of the Performance of the Kernels

In [7]:

```
from sklearn.svm import LinearSVC

stock_df = pd.read_csv("data/Nasdaq.csv", index_col='Date')
print(stock_df.head(3))

daily_return = np.empty(stock_df['Close'].shape)
# From Slides: Daily return (r): Difference in percentage between
# price at time t+1 and time t
daily_return[0] = float('NaN') # The first
daily_return[1:] = np.ediff1d(stock_df['Close']) / stock_df['Close'][:-1]
stock_df.insert(loc=len(stock_df.columns), column='Daily Return', value=daily_return)

binary = (daily_return > 0).astype(float)
stock_df.insert(loc=len(stock_df.columns), column='Binary Decision', value=binary)
stock_df["Binary Decision"] = le.fit_transform(stock_df["Binary Decision"])

stock_df['Rolling Mean 5'] = stock_df['Close'].rolling(5).mean()
stock_df['Rolling Mean 10'] = stock_df['Close'].rolling(10).mean()
stock_df['Rolling Mean 20'] = stock_df['Close'].rolling(20).mean()
stock_df['Rolling Mean 50'] = stock_df['Close'].rolling(50).mean()
stock_df['Rolling Mean 200'] = stock_df['Close'].rolling(200).mean()
stock_df = stock_df.fillna(0) # NAs replaced with zero

stock_models = {
    'Linear SVM Kernel': LinearSVC(dual=False),
    'Poly SVM Kernel': SVC(kernel='poly'),
    'Rbf SVM Kernel': SVC(kernel='rbf'),
    'Sigmoid SVM Kernel': SVC(kernel='sigmoid'),
}
print(stock_df.tail(2))
```

	Open	High	Low	Close	Adj Close	Volume
Date						
1971-02-05	100.000000	100.000000	100.000000	100.000000	100.000000	0
1971-02-08	100.839996	100.839996	100.839996	100.839996	100.839996	0
1971-02-09	100.760002	100.760002	100.760002	100.760002	100.760002	0
	Open	High	Low	Close	\	

Date	Adj Close	Volume	Daily Return	Binary Decision
2021-09-20	14758.139648	14841.820312	14530.070312	14713.900391
2021-09-21	14803.400391	14847.027344	14696.467773	14779.216797

Date	Adj Close	Volume	Daily Return	Binary Decision
2021-09-20	14713.900391	4860630000	-0.021940	0
2021-09-21	14779.216797	3083208000	0.004439	1

Date	Rolling Mean 5	Rolling Mean 10	Rolling Mean 20	Rolling Mean 50
2021-09-20	15027.816016	15126.937012	15143.909961	14875.705195
2021-09-21	14976.107422	15067.425684	15135.738281	14876.624727

Date	Rolling Mean 200
2021-09-20	13856.711787
2021-09-21	13868.721973

With the Rolling Mean, we already have a sort of grouping, thus I vouch to leave the values ungrouped and just work with the Rolling Mean, especially as it is what was also used in prior exercises.

```
In [8]: stock_df["Class Daily Return"] = pd.qcut(stock_df["Daily Return"], q=3)
stock_df["Class Daily Return"], _ = stock_df["Class Daily Return"].factorize(sort=True)

all_stock_features = ['Volume', 'Rolling Mean 5', 'Rolling Mean 10',
                      'Rolling Mean 20', 'Rolling Mean 50', 'Rolling Mean 200']
```

```
In [9]: def kfold_svm_stock_eval(model, x: pd.DataFrame, y: pd.DataFrame):
kf = KFold(n_splits=5)
accuracy = np.empty(kf.n_splits)
precision = np.empty(kf.n_splits)
recall = np.empty(kf.n_splits)
f1 = np.empty(kf.n_splits)

i = 0
for train_index, test_index in kf.split(x):
    x_train, x_test = x.iloc[train_index], x.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    model.fit(x_train, y_train)
    prediction = model.predict(x_test)

    accuracy[i] = accuracy_score(prediction, y_test)
    precision[i] = precision_score(prediction, y_test,
                                   average='macro', zero_division=0)
    recall[i] = recall_score(prediction, y_test,
                              average='macro', zero_division=0)
    f1[i] = f1_score(prediction, y_test,
                     average='macro', zero_division=0)
    i += 1

print(name)
return accuracy, precision, recall, f1
```

```
In [10]: for name, model in stock_models.items():
a, p, r, f = kfold_svm_stock_eval(model, stock_df[all_stock_features],
                                   stock_df["Class Daily Return"])

stock_data = {
    'Fold': [1, 2, 3, 4, 5],
    'Accuracy': a,
    'Precision': p,
    'Recall': r,
    'F1-Score': f,
}

stock_scores = pd.DataFrame(stock_data).set_index('Fold')
display(stock_scores)
```

Linear SVM Kernel

Accuracy	Precision	Recall	F1-Score
----------	-----------	--------	----------

Fold	Accuracy	Precision	Recall	F1-Score
1	0.398227	0.332787	0.216155	0.190712
2	0.304313	0.319782	0.215310	0.227100
3	0.385484	0.333333	0.128495	0.185487
4	0.364516	0.333333	0.121505	0.178093
5	0.317339	0.333333	0.105780	0.160596

Poly SVM Kernel

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.289802	0.333333	0.096601	0.149792
2	0.289399	0.333333	0.096466	0.149630
3	0.256452	0.333333	0.085484	0.136072
4	0.251210	0.333333	0.083770	0.133892
5	0.354435	0.365574	0.264213	0.262435

Rbf SVM Kernel

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.399033	0.333333	0.133011	0.190147
2	0.414349	0.333009	0.138284	0.195419
3	0.310081	0.374814	0.361534	0.238784
4	0.377823	0.341705	0.606177	0.265555
5	0.315726	0.331403	0.297484	0.166609

Sigmoid SVM Kernel

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.289802	0.333333	0.096601	0.149792
2	0.285369	0.323368	0.193278	0.225996
3	0.362097	0.335218	0.282369	0.213611
4	0.368952	0.331721	0.244251	0.272592
5	0.326613	0.325263	0.194355	0.220808

It seems like the poly kernel performs the worst, while rbf performs the best.

Sigmoid is slightly behind linear in performance.