

## Exercise 9 - Ensemble Learning

First name: Brian

Last name: Schweigler

Matriculation number: 16-102-071

(1) Take the titanic dataset and using all attributes to predict the class 'Survived' (convert age and fare into classes ; exclude names from the attribute list) Build a boosting ensemble model with:

- (a) Adaboost
- (b) Gradientboost
- (c) XGB

Show the Comparison of the Performance of the models

```
In [1]: %load_ext autoreload
%autoreload 2
%matplotlib inline

import pandas as pd
import numpy as np
from sklearn import preprocessing
from mlxtend.evaluate import accuracy_score
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import KFold
from sklearn.ensemble import GradientBoostingClassifier, AdaBoostClassifier, VotingClassifier
from xgboost.sklearn import XGBClassifier
from IPython.core.display import display
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier

df = pd.read_csv("data/titanic.csv", index_col='Name')
pd.set_option('display.max_colwidth', None)
le = preprocessing.LabelEncoder()
bins = [0, 4, 18, 65, 100]
labels = ['Infant', 'Child', 'Adult', 'Elderly']
labels = [1, 2, 3, 4]
df['AgeGroup'] = pd.cut(df['Age'], bins=bins, labels=labels, right=False)
df["AgeGroup"] = le.fit_transform(df["AgeGroup"])
df['FareGroup'] = pd.qcut(x=df['Fare'], q=4)
df["FareGroup"] = le.fit_transform(df["FareGroup"])
df["Survived"] = le.fit_transform(df["Survived"])
df["Sex"] = le.fit_transform(df["Sex"])
print(df.head(2))

models = {
    'AdaBoostClassifier': AdaBoostClassifier(),
    'Gradientboost': GradientBoostingClassifier(learning_rate=0.01),
    'XGBClassifier': XGBClassifier(),
}
```

	Survived	Pclass	Sex	\
Name				
Mr. Owen Harris Braund	0	3	1	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	1	1	0	

	Age	\
Name		
Mr. Owen Harris Braund	22.0	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	38.0	

	Siblings/Spouses Aboard	\
Name		
Mr. Owen Harris Braund	1	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	1	

	Parents/Children Aboard	\
Name		
Mr. Owen Harris Braund	0	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	0	

	Fare	AgeGroup	\
Name			
Mr. Owen Harris Braund	7.2500	2	
Mrs. John Bradley (Florence Briggs Thayer) Cumings	71.2833	2	

	FareGroup
Name	

Mr. Owen Harris Braund	0
Mrs. John Bradley (Florence Briggs Thayer) Cumings	3

```
In [2]: all_features = ['Pclass', 'Sex', 'Siblings/Spouses Aboard',
                    'Parents/Children Aboard', 'AgeGroup', 'FareGroup']

def kfold_boost_eval(model, x: pd.DataFrame, y: pd.DataFrame):
    kf = KFold(n_splits=5, shuffle=True, random_state=6)
    accuracy = np.empty(kf.n_splits)
    precision = np.empty(kf.n_splits)
    recall = np.empty(kf.n_splits)
    f1 = np.empty(kf.n_splits)

    i = 0
    for train_index, test_index in kf.split(x):
        x_train, x_test = x.iloc[train_index], x.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        model.fit(x_train, y_train)
        prediction = model.predict(x_test)

        accuracy[i] = accuracy_score(prediction, y_test)
        precision[i] = precision_score(prediction, y_test)
        recall[i] = recall_score(prediction, y_test)
        f1[i] = f1_score(prediction, y_test)
        i += 1

    print(name)
    return accuracy, precision, recall, f1

for name, model in models.items():
    a, p, r, f = kfold_boost_eval(model, df[all_features], df["Survived"])
    data = {
        'Fold': [1, 2, 3, 4, 5],
        'Accuracy': a,
        'Precision': p,
        'Recall': r,
        'F1-Score': f,
    }

    scores = pd.DataFrame(data).set_index('Fold')
    display(scores)
```

AdaBoostClassifier

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.820225	0.779412	0.757143	0.768116
2	0.764045	0.611111	0.758621	0.676923
3	0.779661	0.716418	0.705882	0.711111
4	0.762712	0.657143	0.718750	0.686567
5	0.813559	0.769231	0.735294	0.751880

Gradientboost

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.848315	0.632353	0.955556	0.761062
2	0.786517	0.541667	0.886364	0.672414
3	0.819209	0.656716	0.830189	0.733333
4	0.790960	0.542857	0.883721	0.672566
5	0.847458	0.646154	0.913043	0.756757

XGBClassifier

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.837079	0.691176	0.854545	0.764228
2	0.814607	0.611111	0.897959	0.727273
3	0.790960	0.776119	0.702703	0.737589

	Accuracy	Precision	Recall	F1-Score
<b>Fold</b>				
4	0.768362	0.671429	0.723077	0.696296
5	0.847458	0.769231	0.806452	0.787402

Seems like Gradientboost performs the best out of all 3.

**(2) With your selected stock / market index using all attributes to predict 'daily returns'(decision). ('daily returns' must first be converted into a decision class that will be used as the target(label), all other attributes must be grouped into classes)**

(a) Build a voting ensemble model and Explain how the voting technique affects the performance of the model. (b) Stack any models of your choice to create an ensemble. How does stacking compare with voting?

```
In [3]: from typing import Any

stock_df = pd.read_csv("data/Nasdaq.csv", index_col='Date')
print(stock_df.head(3))

daily_return = np.empty(stock_df['Close'].shape)
# From Slides: Daily return (r): Difference in percentage between
# price at time t+1 and time t
daily_return[0] = float('NaN') # The first
daily_return[1:] = np.ediff1d(stock_df['Close']) / stock_df['Close'][:-1]
stock_df.insert(loc=len(stock_df.columns), column='Daily Return', value=daily_return)

binary = (daily_return > 0).astype(float)
stock_df.insert(loc=len(stock_df.columns), column='Binary Decision', value=binary)
stock_df["Binary Decision"] = le.fit_transform(stock_df["Binary Decision"])

stock_df['Rolling Mean 5'] = stock_df['Close'].rolling(5).mean()
stock_df['Rolling Mean 10'] = stock_df['Close'].rolling(10).mean()
stock_df['Rolling Mean 20'] = stock_df['Close'].rolling(20).mean()
stock_df['Rolling Mean 50'] = stock_df['Close'].rolling(50).mean()
stock_df['Rolling Mean 200'] = stock_df['Close'].rolling(200).mean()
stock_df = stock_df.fillna(0) # NAs replaced with zero

def estimators() -> list[tuple[str, Any]]:
    return [
        ('KNN', KNeighborsClassifier(n_neighbors=3)),
        ('Tree', DecisionTreeClassifier()),
        ('Gaussian Bayes', GaussianNB())
    ]

# Rest is as you'd expect with fit and predict

print(stock_df.tail(2))
```

	Open	High	Low	Close	Adj Close	Volume
Date						
1971-02-05	100.000000	100.000000	100.000000	100.000000	100.000000	0
1971-02-08	100.839996	100.839996	100.839996	100.839996	100.839996	0
1971-02-09	100.760002	100.760002	100.760002	100.760002	100.760002	0

  

	Open	High	Low	Close \
Date				
2021-09-20	14758.139648	14841.820312	14530.070312	14713.900391
2021-09-21	14803.400391	14847.027344	14696.467773	14779.216797

  

	Adj Close	Volume	Daily Return	Binary Decision \
Date				
2021-09-20	14713.900391	4860630000	-0.021940	0
2021-09-21	14779.216797	3083208000	0.004439	1

  

	Rolling Mean 5	Rolling Mean 10	Rolling Mean 20	Rolling Mean 50 \
Date				
2021-09-20	15027.816016	15126.937012	15143.909961	14875.705195
2021-09-21	14976.107422	15067.425684	15135.738281	14876.624727

  

	Rolling Mean 200
Date	
2021-09-20	13856.711787
2021-09-21	13868.721973

With the Rolling Mean, we already have a sort of grouping, thus I vouch to leave the values ungrouped and just work with the Rolling Mean, especially as it is what was also used in prior exercises.

In [4]:

```
stock_df["Class Daily Return"] = pd.qcut(stock_df["Daily Return"], q=3)
stock_df["Class Daily Return"], _ = stock_df["Class Daily Return"].factorize(sort=True)

all_stock_features = ['Volume', 'Rolling Mean 5', 'Rolling Mean 10',
                      'Rolling Mean 20', 'Rolling Mean 50', 'Rolling Mean 200']

def evaluate(model: Any, X_train: pd.DataFrame, X_test: pd.DataFrame, y_train: pd.DataFrame, y_test: pd.DataFrame,
            average: Any = 'binary', zero_division: Any = 'warn') -> (float, float, float, float):
    """
    Evaluates a model on a training and test set.
    :param model: The model to evaluate
    :param average: default='binary': This parameter is required for multiclass/multilabel targets.
    :param zero_division: default='warn'
    :return:
    """
    model.fit(X_train, y_train)
    pred = model.predict(X_test)

    accuracy = accuracy_score(pred, y_test)
    precision = precision_score(pred, y_test, average=average, zero_division=zero_division)
    recall = recall_score(pred, y_test, average=average, zero_division=zero_division)
    f1 = f1_score(pred, y_test, average=average, zero_division=zero_division)

    return accuracy, precision, recall, f1

def kfold_eval(model: Any, X: pd.DataFrame, y: pd.DataFrame, k: int, average: Any = 'binary',
               zero_division: Any = 'warn') -> (np.array, np.array, np.array, np.array):
    """
    Performs k-fold cross-validation on the given model.
    :param model: The model to evaluate
    :param X: the features to train on
    :param y: the labels to predict
    :param k: how many folds to perform
    :param average: default='binary': This parameter is required for multiclass/multilabel targets.
    :param zero_division: default='warn'
    :return: accuracy, precision, recall, f1
    """

    kf = KFold(n_splits=k)

    accuracy = np.empty(kf.n_splits)
    precision = np.empty(kf.n_splits)
    recall = np.empty(kf.n_splits)
    f1 = np.empty(kf.n_splits)

    i = 0
    for train_index, test_index in kf.split(X):
        accuracy[i], precision[i], recall[i], f1[i] = evaluate(
            model,
            X.iloc[train_index],
            X.iloc[test_index],
            y.iloc[train_index],
            y.iloc[test_index],
            average,
            zero_division,
        )
        i += 1

    return accuracy, precision, recall, f1

def boosting_performance(model: Any, target_feature: str = 'Daily Returns Class', k: int = 5, average: Any = 'binary',
                        zero_division: Any = 'warn') -> (np.array, np.array, np.array, np.array):
    """
    Performs a k-fold cross-validation on an SVM with the specified kernel.

    :param model: The SVM model
    :param target_feature: default='Daily Returns Class': The target feature
    :param k: how many folds to perform
    :param average: default='binary': This parameter is required for multiclass/multilabel targets.
    :param zero_division: default='warn'

    :return: accuracy, precision, recall, f1
    """
    X = stock_df[all_stock_features]
    y = stock_df["Class Daily Return"]

    return kfold_eval(model=model, X=X, y=y, k=k, average=average, zero_division=zero_division)
```

```
In [5]: k_folds = 5
for voting in ['soft', 'hard']:
    print(f'Voting = {voting}:')
    vote_model = VotingClassifier(estimators=estimators(), voting=voting)

    a, p, r, f = boosting_performance(vote_model, k=k_folds, average='macro', zero_division=0)
    data = {
        'Fold': range(1, k_folds + 1),
        'Accuracy': a,
        'Precision': p,
        'Recall': r,
        'F1-Score': f,
    }

    vote_scores = pd.DataFrame(data).set_index('Fold')
    display(vote_scores)
```

Voting = soft:

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.397420	0.341186	0.334484	0.243252
2	0.313180	0.336974	0.348685	0.254276
3	0.333871	0.313802	0.317459	0.307042
4	0.340726	0.313650	0.308428	0.302055
5	0.317339	0.327554	0.325754	0.261987

Voting = hard:

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.393793	0.334886	0.349451	0.229639
2	0.350262	0.330278	0.358524	0.286219
3	0.310081	0.323959	0.321672	0.305869
4	0.358468	0.329020	0.329986	0.314219
5	0.319355	0.333172	0.403910	0.210987

Hard voting entails picking the prediction with the highest number of votes, whereas soft voting entails combining the probabilities of each prediction in each model and picking the prediction with the highest total probability.

Performance is quite similar in this case here though.

**(b) Stack any models of your choice to create an ensemble. How does stacking compare with voting?**

```
In [6]: from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import StackingClassifier

k_folds = 5
stack_model = StackingClassifier(estimators=estimators(), final_estimator=LogisticRegression(solver='saga', max_iter=5000))

a, p, r, f = boosting_performance(stack_model, k=k_folds, average='macro', zero_division=0)
data = {
    'Fold': range(1, k_folds+1),
    'Accuracy': a,
    'Precision': p,
    'Recall': r,
    'F1-Score': f,
}

stack_scores = pd.DataFrame(data).set_index('Fold')
stack_scores
```

```
Out[6]:
```

	Accuracy	Precision	Recall	F1-Score
Fold				
1	0.312374	0.330195	0.317820	0.204012
2	0.391778	0.337808	0.361225	0.278679
3	0.293952	0.350124	0.353973	0.245671

	Accuracy	Precision	Recall	F1-Score
Fold				
4	0.278226	0.342131	0.353248	0.220650
5	0.372984	0.364003	0.368794	0.295430

The stacking classifier is able to learn when our base estimators can be trusted or not. Stacking allows us to use the strength of each individual estimator by using their output as an input of a final estimator.

Performance seems to be slightly better with voting than with stacking in our case.